

Homework 1

10-605/805: Machine Learning with Large Datasets

Due Monday, September 11th at 11:59PM Eastern Time

Submit your solutions via Gradescope, **with your solution to each subproblem on a separate page**, i.e., following the template below.

IMPORTANT: Be sure to highlight where your solutions are for each question when submitting to Gradescope otherwise you will be marked 0 and will need to submit regrade request for every solution un-highlighted in order for fix it!

Note that Homework 1 consists of two parts: this written assignment, and a programming assignment. The written part and a programming part. Remember to fill out the collaboration section found at the end of this homework as per the course policy.

You are required to complete sections 1.1, 1.2 and the two parts of the programming. Your homework score will be calculated as a percentage over the maximum points you can earn, which is 100.

1 Written Section [50 Points]

1.1 Sparse PCA

Sparse PCA is a variant of PCA that can be used to better interpret which of the k features are most important for the various principal components. In sparse PCA, the first principal component \mathbf{p} is the solution to the following constrained optimization problem:

$$\begin{aligned} \max_{\mathbf{p}} \quad & \mathbf{p}^T \mathbf{C}_X \mathbf{p} \\ \text{subject to} \quad & \|\mathbf{p}\|_2 = 1 \text{ and} \\ & \|\mathbf{p}\|_0 \leq s \end{aligned} \tag{1}$$

Here, \mathbf{X} is an $n \times k$ matrix of n data points and \mathbf{C}_X is the covariance matrix of \mathbf{X} . The L0 norm $\|\mathbf{p}\|_0$ represents the number of non-zero elements in the vector \mathbf{p} , and we want this to be at most s , where s is a fixed constant such that $1 \leq s \leq k$.

- (a) [10 points] The second constraint of the above optimization problem, $\|\mathbf{p}\|_0 \leq s$, is not convex. Show this by proving that the L0 norm $\|\cdot\|_0$ is not a convex function.

Hint - Use Jensen's Inequality

Use Jensen's Inequality, for a convex function f , if X is a random variable, then:

$$f(\mathbb{E}[X]) \leq \mathbb{E}[f(X)] \tag{2}$$

Consider L0 norm, and define two vectors a and b , such that:

$$\begin{aligned} a &= [1, 0, 0] \\ b &= [0, 1, 0] \end{aligned} \tag{3}$$

Then average of a and b :

$$c = \frac{1}{2}(a + b) = \left[\frac{1}{2}, \frac{1}{2}, 0\right] \tag{4}$$

Then the L0 norm of c :

$$\|c\|_0 = 2 \tag{5}$$

While the average of L0 norm of a and L0 norm of b :

$$\frac{1}{2}(\|a\|_0 + \|b\|_0) = \frac{1}{2}(1 + 1) = 1 \tag{6}$$

Obviously, $\|c\|_0 = 2$ is larger than the average of L0 norm of a and L0 norm of b which is 1, this is against the definition of convex function according to equation (2), so L0 norm is not a convex function.

- (b) [10 points] Being a non-convex problem, (1) is hard to solve directly. Instead, we can use the following brute-force optimization method. Fix a set $S \subset \{1, 2, \dots, k\}$, such that $|S| = s$ corresponding to s locations of non-zero elements in \mathbf{p} . For each S we solve the following convex problem:

$$\begin{aligned} \max_{\mathbf{p}} \quad & \mathbf{p}^T \mathbf{C}_X \mathbf{p} \\ \text{subject to} \quad & \|\mathbf{p}\|_2 = 1 \text{ and } p_i = 0 \text{ for } i \notin S, \end{aligned} \quad (7)$$

where p_i is the i -th element of \mathbf{p} . To solve the original problem (1), we can solve (7) for all possible sets S and choose the solution that gives the maximum objective value.

How many different sets S will we have to iterate over in the above method? Why is this brute-force method computationally expensive for high-dimensional data?

To select a set $S \subset \{1, 2, \dots, k\}$, such that $|S| = s$, this is a combination problem, so the number can be denoted as:

$$C(k, s) = \frac{k!}{s! (k-s)!}$$

And s can be chose from 1 to k , so the total is:

$$\sum_{s=1}^k C(k, s) = \sum_{s=1}^k \frac{k!}{s! (k-s)!}$$

So the brute-force method is obviously computationally impractical, because it grows exponentially with the size of K

So for high-dimensional data, the K is very large, it is definitely infeasible to brute-force all the possible situation.

- (c) [5 points] Instead of the brute-force method suggested in (b) which yields an exact solution, we may want to consider approximate solutions that do not necessarily satisfy $\|p\|_0 \leq s$, but impose some sparsity in p . Do you have any suggestions for an approximate solution which can be more computationally tractable? Listing any one method will suffice.

One possible solution is L_1 regularization.
So adding the L_1 regularization makes the original problem become:

$$\begin{aligned} \max_p \quad & p^T C x p - \lambda \|p\|_1 \\ \text{subject to} \quad & \|p\|_2 = 1 \end{aligned}$$

The L_1 penalty term, $\|p\|_1$, will promote the sparsity by encouraging some elements of p to be zero. This can help select a subset of the original features and at the same time, this is a computationally tractable solution.

1.2 Laplacian Eigenmaps

Several non-linear methods for dimensionality reduction assume that high-dimensional data lie on or near a low-dimensional non-linear manifold embedded in the input space. These so-called manifold learning techniques aim to learn this manifold structure. Global methods, e.g., Isomap, aim to preserve distances along the manifold between all pairs of data points. In contrast, local methods, like **Laplacian Eigenmaps** focus only on preserving local neighborhood relationships in the high-dimensional space.

In particular, given a dataset of n data points, Laplacian Eigenmaps works by first finding t nearest neighbors for each point; then constructing \mathbf{W} , the sparse, symmetric $n \times n$ nearest neighborhood adjacency matrix, where $\mathbf{W}_{ij} = \exp(-\|\mathbf{x}_i - \mathbf{x}_j\|_2^2 / \sigma^2)$ if $(\mathbf{x}_i, \mathbf{x}_j)$ are neighbors¹, 0 otherwise, and σ is a scaling parameter; and finally, finding the r -dimensional representation that minimizes the following objective:

$$\mathbf{Y} = \operatorname{argmin}_{\mathbf{Y}'} \sum_{i,j} \mathbf{W}_{ij} \|\mathbf{y}'_i - \mathbf{y}'_j\|_2^2, \quad (8)$$

subject to appropriate constraints on \mathbf{Y}' . Note that $\mathbf{Y}' \in \mathbb{R}^{r \times n}$ and $\mathbf{y}'_i \in \mathbb{R}^r$ is the i th column of \mathbf{Y}' .

- (a) **[5 points]** Describe in words how the objective in Eq (8) aims to preserve local neighborhood relationships. What is the role of the weight matrix \mathbf{W} ?

The objective in Eq(8) aims to preserve local neighborhood relationships by imposing $n \times n$ nearest neighborhood matrix \mathbf{W} , where

$$\mathbf{W}_{ij} = \exp\left(-\|\mathbf{x}_i - \mathbf{x}_j\|_2^2 / \sigma^2\right), \text{ if } (\mathbf{x}_i, \mathbf{x}_j) \text{ is neighbor}$$

This matrix implies that if $\mathbf{x}_i, \mathbf{x}_j$ is neighbor then \mathbf{W}_{ij} will have a positive value, which means the corresponding $\|\mathbf{y}'_j - \mathbf{y}'_i\|_2^2$ will contribute significantly to the objective.

¹To ensure that \mathbf{W} is symmetric, we define $(\mathbf{x}_i, \mathbf{x}_j)$ as neighbors if either data point is a t nearest neighbor of the other.

Whereas, if x_i, x_j is not neighbor, then w_{ij} would correspond to zero, and the corresponding $\|y_j' - y_i'\|_2^2$ would have no contribution to the objective function, which means it will not make y_j' and y_i' to be near in the low dimension.

Also, we can see that as $\|x_i - x_j\|_2^2$ increases, w_{ij} will decrease, which means that the larger the distance between x_i and x_j , the less $\|y_j' - y_i'\|_2^2$ is contributed to the objective function, which means y_j' and y_i' will be further away in the low-dimension space.

- (b) [10 points] Assume we seek a one-dimensional representation \mathbf{y} , i.e., $r = 1$. Show that (8) is equivalent to

$$\mathbf{y} = \operatorname{argmin}_{\mathbf{y}'} \mathbf{y}'^\top \mathbf{L} \mathbf{y}', \quad (9)$$

where \mathbf{D} is the diagonal matrix with $D_{ii} = \sum_j W_{ij}$, and $\mathbf{L} = \mathbf{D} - \mathbf{W}$ is the graph Laplacian.

Hint - Expand (8) and use the distributive property of summation

Expand equation (8) :

$$\begin{aligned} Y &= \operatorname{argmin}_{\mathbf{y}'} \sum_{i,j} W_{i,j} \| \mathbf{y}'_j - \mathbf{y}'_i \|_2^2 \\ &= \operatorname{argmin}_{\mathbf{y}'} \sum_{i,j} W_{i,j} \| \mathbf{y}'_j - \mathbf{y}'_i \|_2^2 \quad (r=1) \\ &= \operatorname{argmin}_{\mathbf{y}'} \sum_{i,j} W_{i,j} (\mathbf{y}'_j - \mathbf{y}'_i)^2 \\ &= \operatorname{argmin}_{\mathbf{y}'} \sum_{i,j} W_{i,j} (\mathbf{y}'_j{}^2 - 2\mathbf{y}'_j \mathbf{y}'_i + \mathbf{y}'_i{}^2) \\ &= \operatorname{argmin}_{\mathbf{y}'} \sum_{i,j} W_{i,j} \mathbf{y}'_j{}^2 - 2 \sum_{i,j} W_{i,j} \mathbf{y}'_j \mathbf{y}'_i \\ &\quad + \sum_{i,j} W_{i,j} \mathbf{y}'_i{}^2 \end{aligned}$$

Use the distributive property of summation.

This will become :

$$\underset{y'}{\operatorname{argmin}} \sum_j y_j'^2 \sum_i w_{ij} - 2 \sum_{i,j} w_{ij} y_j' y_i'$$

$$+ \sum_i y_i'^2 \sum_j w_{ij}$$

And by definition, $\sum_j y_j'^2 \sum_i w_{ij} = y'^T D y'$

$$\text{also } \sum_i y_i'^2 \sum_j w_{ij} = y'^T D y'$$

So it becomes :

$$\underset{y'}{\operatorname{argmin}} 2 \left(y'^T D y' - 2 \sum_{i,j} w_{ij} y_j' y_i' \right)$$

$$\text{Also by definition, } \sum_{i,j} w_{ij} y_j' y_i' = y'^T W y'$$

So it becomes :

$$\underset{y'}{\operatorname{argmin}} 2 \left(y'^T D y' - y'^T W y' \right)$$

$$= \underset{y'}{\operatorname{argmin}} y'^T (D - W) y'$$

$$= \underset{y'}{\operatorname{argmin}} y'^T L y'$$

- (c) [5 points] From (8) and (9), we can see that the graph Laplacian is positive semidefinite (take a minute to figure out why this is). The solution to the problem is anything proportional to the ones vector. Explain why this is the solution and why it might be bad.

Because L is semidefinite, which means

$$y'^T L y' \geq 0, \quad \text{for all } y' \in \mathbb{R}^r$$

And according to the property of semidefinite, the sum of every row equals to zero.

Thus, if y' is ones vector, then $L y' = 0$, which means $y'^T L y' = 0$, if $y' \propto$ ones vector

And we know $y'^T L y' \geq 0$, which means y' proportional to ones vector is the solution to the problem.

The solution of y' which proportional to ones means basically all the elements of low-dim representation is the same.

The could be bad for capturing local structure, might losing local information.

(d) [5 points] What constraint can you put on \mathbf{y}' to stop the ones vector from being the solution?

One possible constraint could be :

$$\sum_{i=1}^r y_i' = 0$$

This can force the objective function to avoid from ones vector solution,

2 Programming [50 points]

This assignment involves understanding some basics of distributed computing, the MapReduce programming model, Spark, and an example of data cleaning.

This assignment consists of two major parts. The first part is a tutorial on entity resolution, a common type of data cleaning, and the second part is a walk-through of a typical machine learning pipeline on a real-world dataset.

We provide the code template for this assignment in *two* Jupyter notebooks. What you need to do is to follow the instructions in the notebooks and implement the missing parts marked with '`<FILL IN>`' or '`# YOUR CODE HERE`'. Most of the '`<FILL IN>/YOUR CODE HERE`' sections can be implemented in just one or two lines of code. Also, keep in mind to delete the '`raise NotImplementedError()`' once you are done implementing a function. We provide several `assert` statements in the notebook for you to check the validity of your solution.

2.1 Getting lab files

You can find the notebooks '`hw1_coding_1.ipynb`' and '`hw1_coding_2.ipynb`' in the homework 1 handout .tar file.

Next, import the notebooks into your Databricks account, which provides you a well-configured Spark environment and will definitely save your time (see the next section for details).

2.2 Setting up environments on Databricks

We provide step-by-step instructions on how to configure your Databricks platform. We will also introduce it in detail during the recitation on September 1st. Here are the links to the recitation 1 materials for your convenience: [Recitations 1 Slides](#), [Lab Notebook](#).

1. Sign up for the **Community Edition** of Databricks here: <https://databricks.com/try-databricks>.
2. Import the notebook file we provide on your homepage: **Workspace -> Users -> Import**
3. Create a cluster: **Clusters -> Create Cluster**. You can use any cluster name as you like. When configuring your cluster, make sure to choose **runtime version 13.2**. Note: It may take a while to launch the cluster, please wait for its status to turn to '**active**' before start running.
4. Installing third-party packages that will be used in the homework on Databricks: **Clusters -> Cluster name -> Libraries -> Install New**. Then select PyPI, enter the package name as **nose**. Finally click **Install** to install it.
5. You can start to play with the notebook now!

Note: Databricks Community Edition only allows you to launch one 'cluster'. If the current cluster is 'terminated', then you can either (1) delete it, and then create a new one, or (2) activate and attach to the existing cluster when running the notebook. Make sure to install nose.

2.3 Entity Resolution

Entity Resolution, or "Record linkage" is the term used by statisticians, epidemiologists, and historians, among others, to describe the process of joining records from one data source with another that describe the same entity. Other terms with the same meaning include, "entity disambiguation/linking", "duplicate detection", "deduplication", "record matching", "(reference) reconciliation", "object identification", "data/information integration", and "conflation".

Entity Resolution (ER) refers to the task of finding records in a dataset that refer to the same entity across different data sources (e.g., data files, books, websites, databases). ER is necessary when joining datasets based on entities that may or may not share a common identifier (e.g., database key, URI, National identification number), as the case may be due to differences in record shape, storage location, and/or curator style or preference. A dataset that has undergone ER may be referred to as being cross-linked. In this homework, we break the entity resolution problem into four parts:

- **Part 0 – Preliminaries:** Load in the dataset into pair RDDs where the key is the mapping ID, and the value is a string consisting of the name/title, description, and manufacturer of the corresponding record.
- **Part 1 – ER as Text Similarity - Bags of Words:** Build components for bag-of-words text analysis, and then compute record similarity. Bag-of-words is a conceptually simple yet powerful approach for text analysis. The idea is treating strings, a.k.a. documents, as unordered collections of words or tokens, i.e., as bags of words.
- **Part 2 – ER as Text Similarity - Weighted Bag-of-Words using TF-IDF:** In this part we compute the TF-IDF weight for each record. Bag-of-words comparisons do not perform well when all tokens are treated in the same way. In real world scenarios, some tokens are more important than the others. Weights give us a way to specify which tokens could have higher "importance". With weights, when we compare documents, instead of counting common tokens, we sum up the weights of common tokens. A good heuristic for assigning weights is called "Term-Frequency/Inverse-Document-Frequency," or TF-IDF for short. **TF** rewards tokens that appear many times in the same document. It is computed as the frequency of a token in a document. **IDF** rewards tokens that are rare overall in a dataset. The intuition is that it is more significant if two documents share a rare word than a common one.
- **Part 3 – ER as Text Similarity - Cosine Similarity:** Compute the cosine similarity of the tokenized strings based on the TF-IDF weights.
- **Part 4 – Scalable ER:** Use the inverted index data structure to scale ER. The ER algorithm above is quadratic in two ways. First, we did a lot of redundant computation of tokens and weights, since each record was reprocessed every time it was compared. Second, we made quadratically many token comparisons between records. In reality, most records have nothing (or very little) in common. Moreover, it is typical for a record in one dataset to have at most one duplicate record in the other dataset (this is the case assuming each dataset has been de-duplicated against itself). In this case, the output is linear in the size of the input and we can hope to achieve linear running time. An inverted index is a data structure that will allow us to avoid making quadratically many token comparisons. It maps each token in the dataset to the list of documents that contain the token. So, instead of comparing, record by record, each token to every other token to see if they match, we will use inverted indices to look up records that match on a particular token.
- **Part 5 – Analysis:** Determine duplicate entities based on the similarity scores, and compute evaluation metrics. Now we have an authoritative list of record-pair similarities, but we need a way to

use those similarities to decide if two records are duplicates or not. The simplest approach is to pick a threshold. Different thresholds correspond to different false positives and false negatives, which will result in different precision and recall scores.

2.4 Machine Learning Pipeline

Following what we have talked about in class, we will work through a minimal machine learning pipeline. The steps include understanding the problem to obtaining a machine learning model. We break this down into 7 parts.

- **Part 1 – Business Understanding:** The first step is to understand the problem we are trying to solve. This could involve conversations with stakeholders as well as readings describing the problem at hand.
- **Part 2 – Data Extract-Transform-Load:** We need to take the raw data we are given and make it usable for machine learning purposes, potentially pulling it from multiple sources.
- **Part 3 – Explore Your Data:** In this part we use SQL embedded in notebooks to take a look at the underlying database containing the data.
- **Part 4 – Visualize Your Data:** In this section, we use the notebook to create a few visualizations of the data we have.
- **Part 5 – Data Preparation:** We have to group up the relevant variables we want to actually use in our model.
- **Part 6 – Data Modeling:** Create a machine learning model and train it.
- **Part 7 – Tuning and Evaluation:** Try out different hyperparameter settings and different model types.

See the notebooks for detailed descriptions and instructions of each question.

3 Collaboration Questions

1. (a) Did you receive any help whatsoever from anyone in solving this assignment?

(b) If you answered 'yes', give full details (e.g. "Jane Doe explained to me what is asked in Question 3.4")

(a) No

2. (a) Did you give any help whatsoever to anyone in solving this assignment?

(b) If you answered 'yes', give full details (e.g. "I pointed Joe Smith to section 2.3 since he didn't know how to proceed with Question 2")

(a) No

3. (a) Did you find or come across code that implements any part of this assignment?

(b) If you answered 'yes', give full details (book & page, URL & location within the page, etc.).

(a) No